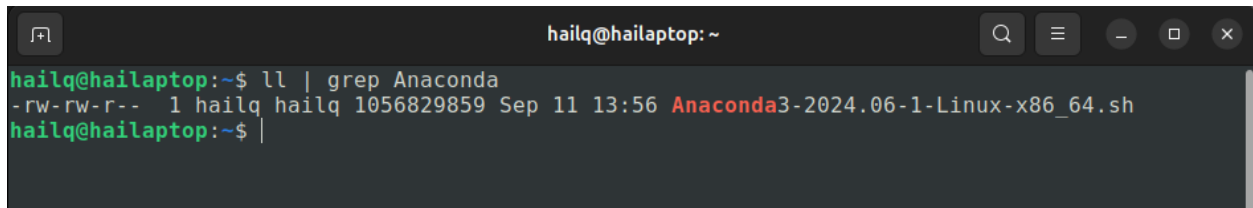


Tutorial 1. Getting Started

Getting started with Anaconda

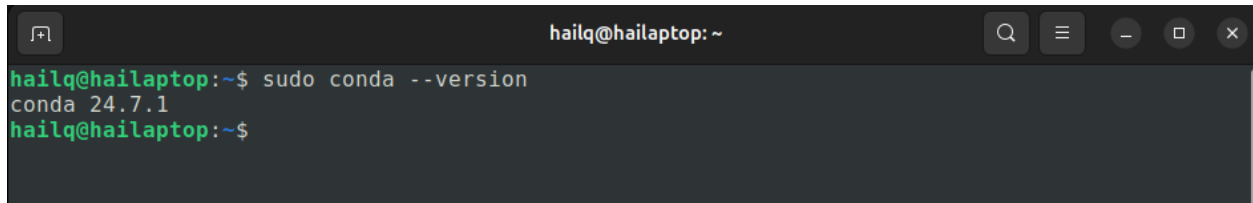
Install conda

1. Downloaded the installer package from the website. The screenshot shows there is an executable file for installing conda in the directory.



```
hailq@hailaptop: ~  
hailq@hailaptop:~$ ll | grep Anaconda  
-rw-rw-r-- 1 hailq hailq 1056829859 Sep 11 13:56 Anaconda3-2024.06-1-Linux-x86_64.sh  
hailq@hailaptop:~$
```

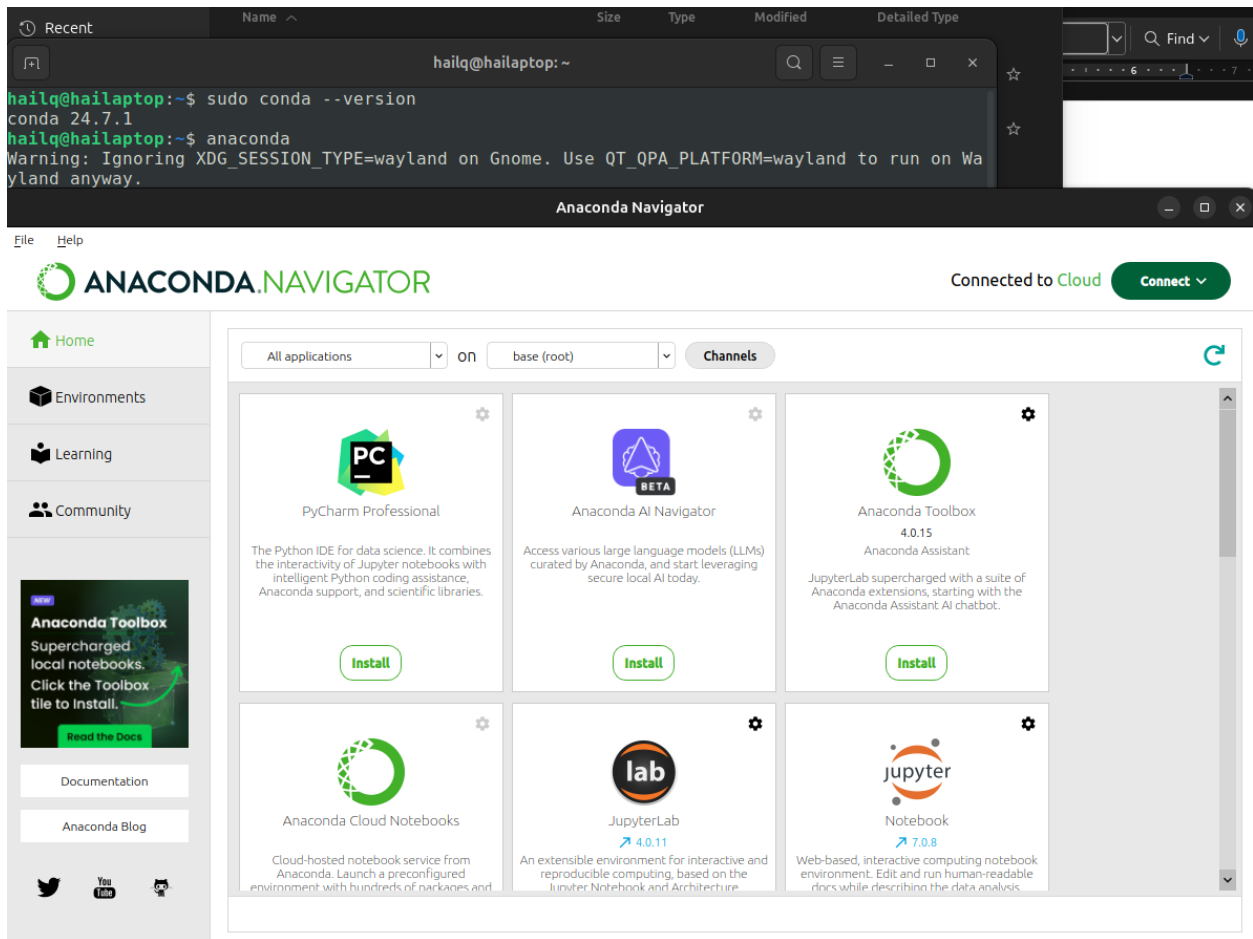
2. Execute the installer, the message below verifies that conda version 24.7.1 has been successfully installed.



```
hailq@hailaptop: ~  
hailq@hailaptop:~$ sudo conda --version  
conda 24.7.1  
hailq@hailaptop:~$
```

Install anaconda

1. Install anaconda-navigator using conda. Having configured a symlink in environment variable to the /anaconda/bin/anaconda-navigator, open the anaconda-navigator via terminal. The screenshot below verifies successful installation



Install jupyter packages via conda

1. Install jupyter via terminal (ubuntu)

```
hailq@hailaptop:~$ conda install jupyter
Channels:
 - defaults
Platform: linux-64
```

2. Create symlink to the jupyter-notebook file in /usr/local/bin so that jupyter notebook can be opened via terminal (ubuntu)

```
hailq@hailaptop:~$ ll /usr/local/bin/ | grep jupyter
lrwxrwxrwx 1 root root 35 Sep 12 10:31 jupyter -> /usr/local/src/anaconda/bin/jupyter*
lrwxrwxrwx 1 root root 44 Sep 12 10:33 jupyter-notebook -> /usr/local/src/anaconda/bin/jupyter-notebook*
```

Run Jupyter Notebook

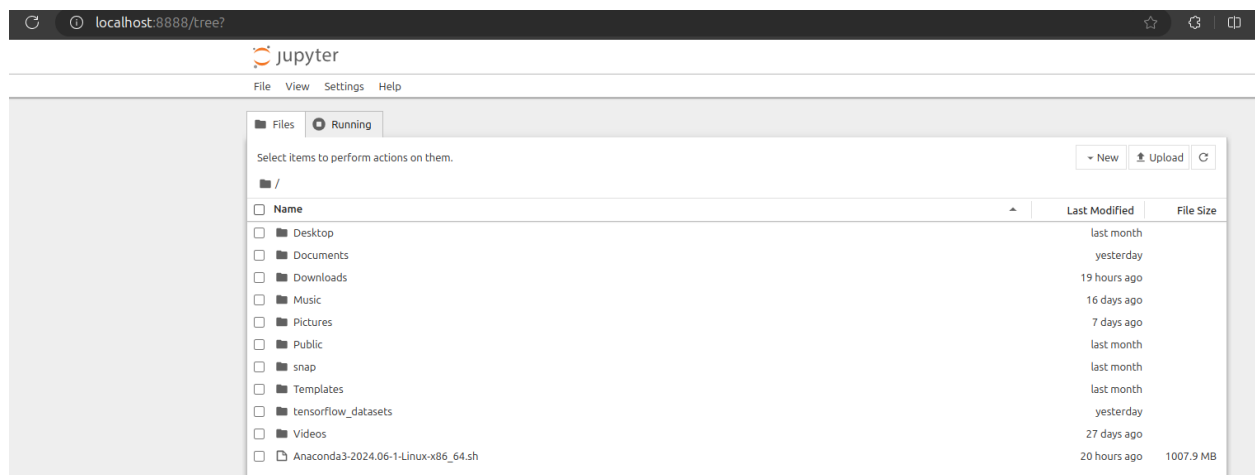
1. Execute jupyter-notebook package

```

hailq@hailaptop:~$ jupyter notebook
[W 2024-09-12 10:53:30.315 ServerApp] A `_jupyter_server_extension_points` function was not found in jupyter_lsp. Instead, a `_jupyter_server_extension_paths` function was found and will be used for now. This function name will be deprecated in future releases of Jupyter Server.
[W 2024-09-12 10:53:30.640 ServerApp] A `_jupyter_server_extension_points` function was not found in notebook_shim. Instead, a `_jupyter_server_extension_paths` function was found and will be used for now. This function name will be deprecated in future releases of Jupyter Server.
[I 2024-09-12 10:53:31.384 ServerApp] Extension package panel.io.jupyter_server_extension took 0.7427s to import
[I 2024-09-12 10:53:31.384 ServerApp] jupyter_lsp | extension was successfully linked.
[I 2024-09-12 10:53:31.389 ServerApp] jupyter_server_terminals | extension was successfully linked.
[I 2024-09-12 10:53:31.394 ServerApp] jupyterlab | extension was successfully linked.
[I 2024-09-12 10:53:31.398 ServerApp] notebook | extension was successfully linked.
[I 2024-09-12 10:53:31.617 ServerApp] notebook_shim | extension was successfully linked.
[I 2024-09-12 10:53:31.617 ServerApp] panel.io.jupyter_server_extension | extension was successfully linked.

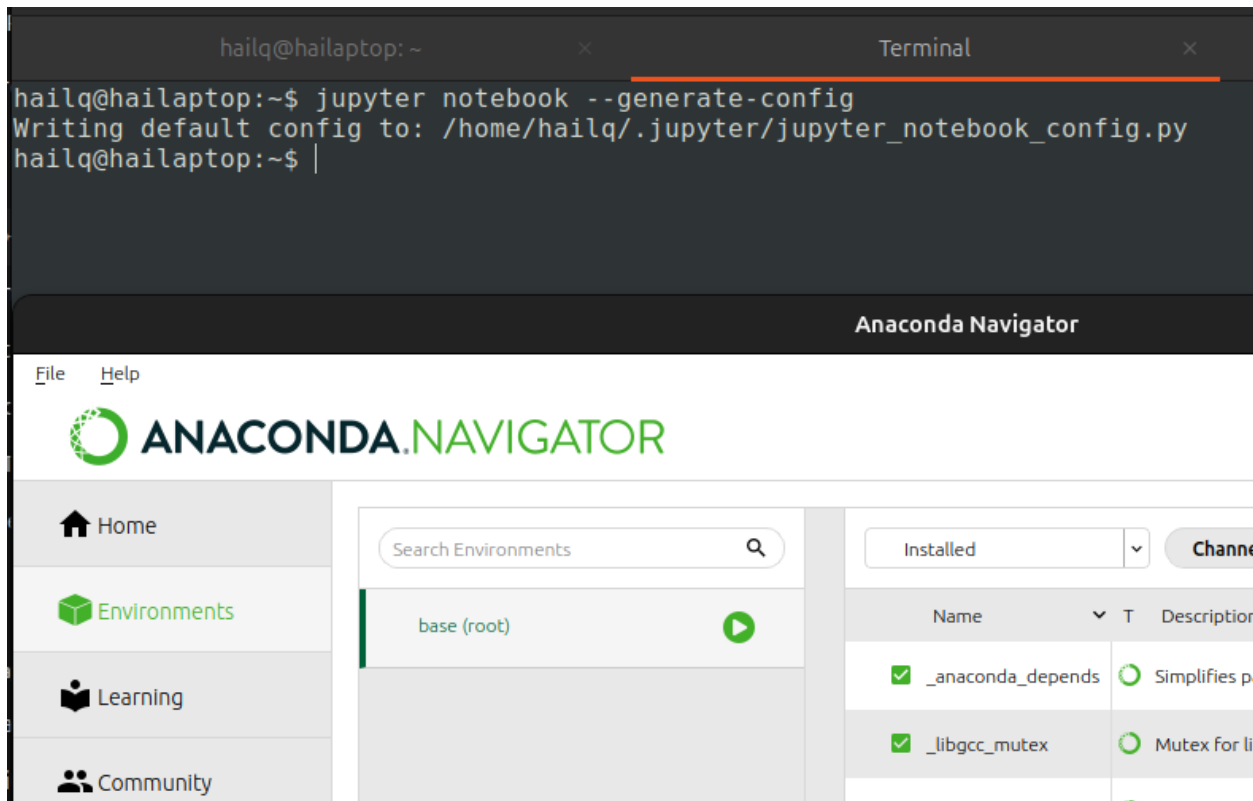
```

2. Enter jupyter notebook GUI via 127.0.0.1:8888 endpoint

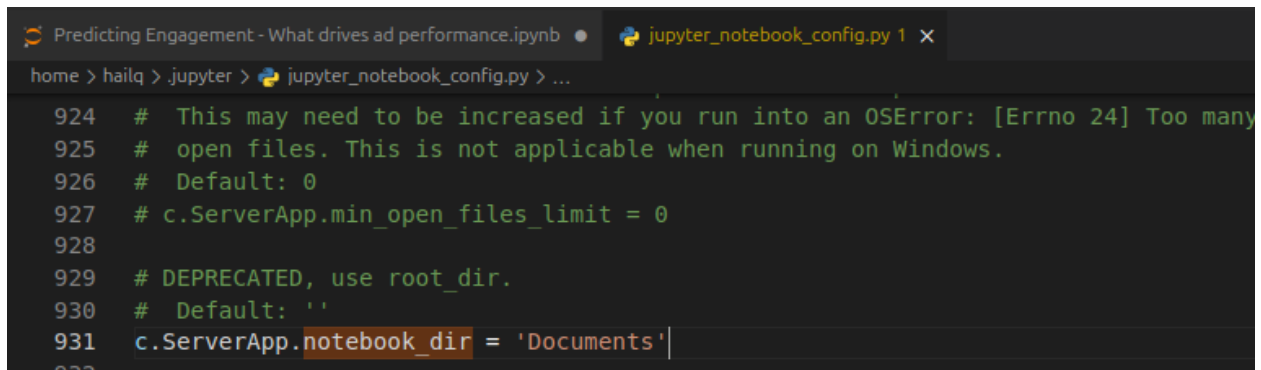


Change Jupyter Notebook start directory

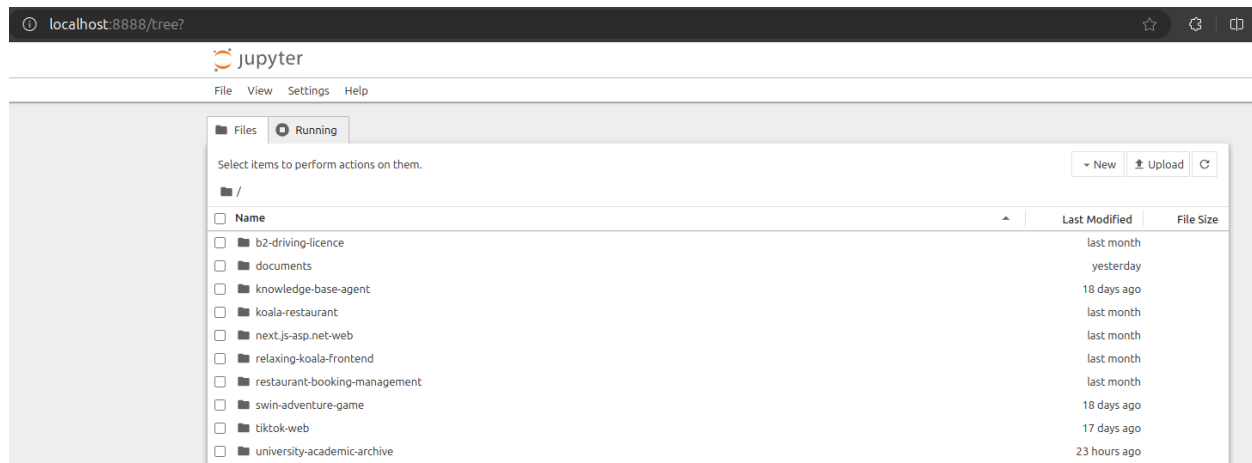
1. Open the Anaconda Navigator and click on Environments -> base(root) -> Open Terminal. Type the command `jupyter notebook --generate-config`. On ubuntu system, a configuration file is created at `~` which is `/home/username/.jupyter/jupyter_notebook_config.py`



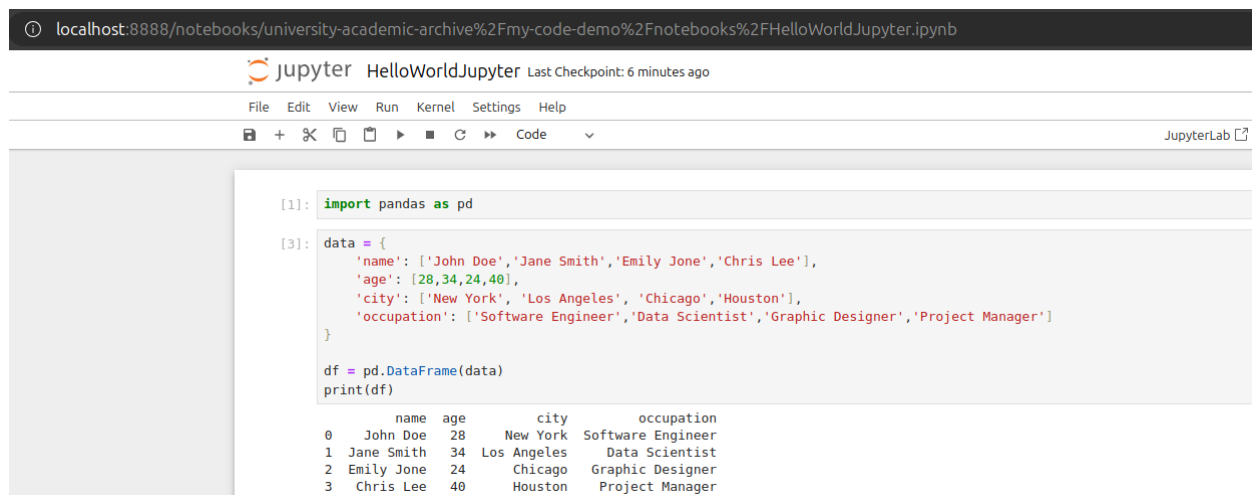
2. Modify the configuration in the `jupyter_notebook_config.py` to `~/Documents`



3. Verify the new starting directory applied to jupyter notebook. We no longer access home directory by default (that contains Documents, Download, Pictures, ...)



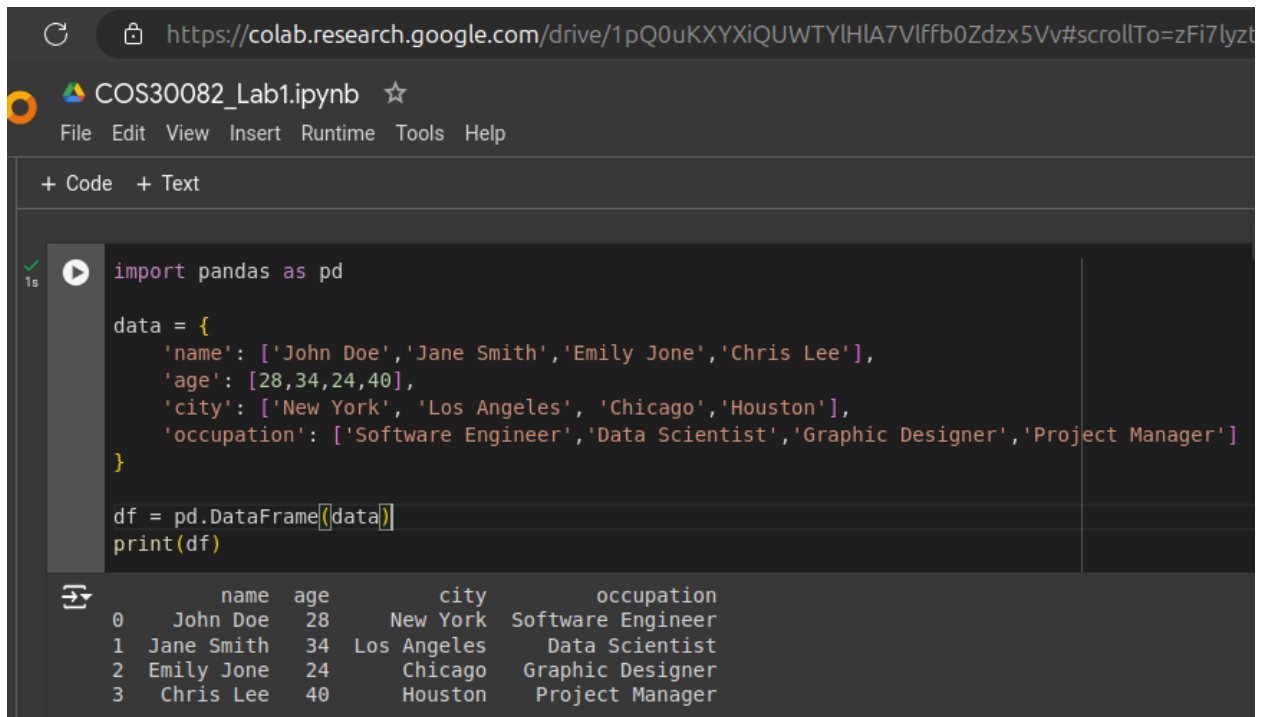
Run some Python code



Getting started with Google Collab

Familiarize with the interface

1. Login to Google Collab, create a new Notebook, and write some code then execute the code



The screenshot shows a Google Colab notebook interface. The browser address bar displays the URL: `https://colab.research.google.com/drive/1pQ0uKXYXiQUWTYIHA7Vlffb0Zdx5Vv#scrollTo=zFi7lyzt`. The notebook title is `COS30082_Lab1.ipynb`. The menu bar includes `File`, `Edit`, `View`, `Insert`, `Runtime`, `Tools`, and `Help`. Below the menu, there are tabs for `+ Code` and `+ Text`. The code cell contains the following Python code:

```
import pandas as pd

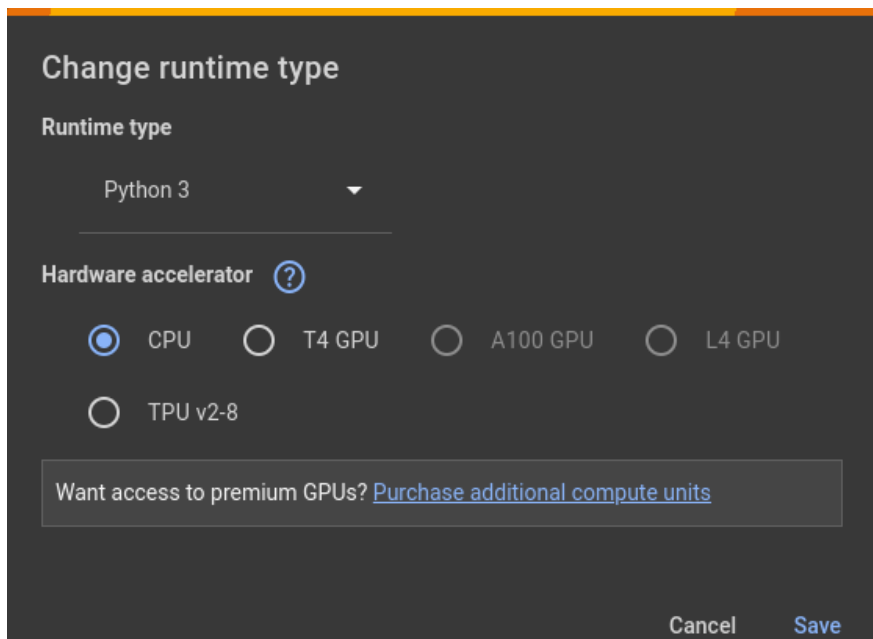
data = {
    'name': ['John Doe', 'Jane Smith', 'Emily Jone', 'Chris Lee'],
    'age': [28, 34, 24, 40],
    'city': ['New York', 'Los Angeles', 'Chicago', 'Houston'],
    'occupation': ['Software Engineer', 'Data Scientist', 'Graphic Designer', 'Project Manager']
}

df = pd.DataFrame(data)
print(df)
```

The output of the code is a pandas DataFrame with 4 rows and 5 columns:

| | name | age | city | occupation |
|---|------------|-----|-------------|-------------------|
| 0 | John Doe | 28 | New York | Software Engineer |
| 1 | Jane Smith | 34 | Los Angeles | Data Scientist |
| 2 | Emily Jone | 24 | Chicago | Graphic Designer |
| 3 | Chris Lee | 40 | Houston | Project Manager |

Accessing GPU and TPU resources

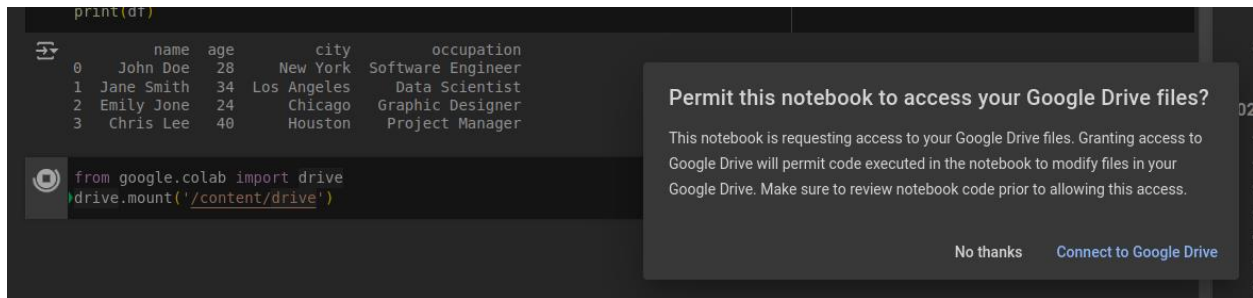


The screenshot shows the 'Change runtime type' dialog in Google Colab. The 'Runtime type' is set to 'Python 3'. The 'Hardware accelerator' section has a question mark icon. Below it, there are five radio button options: `CPU` (selected), `T4 GPU`, `A100 GPU`, `L4 GPU`, and `TPU v2-8`. At the bottom, there is a text box that says: 'Want access to premium GPUs? [Purchase additional compute units](#)'. The dialog has 'Cancel' and 'Save' buttons at the bottom right.

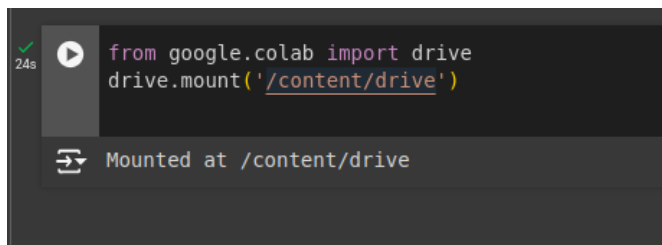
Mount Google Drive

1. Execute the code as instructed

```
from google.colab import drive
drive.mount('/content/drive')
```



2. Verify credentials and authenticate. Verify mounted successfully

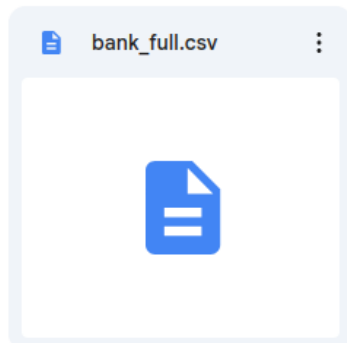


3. Now when upload a file to /Collab Notebook/datasets

Drive của tôi > Colab Notebooks > datasets ▾

Loại ▾ Người ▾ Lần sửa đổi gần đây nhất ▾

Tệp



4. The dataset can be accessed via pandas

```
✓ 22s [3] from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

✓ 0s import pandas as pd

df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/datasets/bank_full.csv')
print(df)
```

| | age | job | marital | education | default | balance | housing | loan | \ |
|-------|-----|--------------|----------|-----------|---------|---------|---------|------|---|
| 0 | 58 | management | married | tertiary | no | 2143 | yes | no | |
| 1 | 44 | technician | single | secondary | no | 29 | yes | no | |
| 2 | 33 | entrepreneur | married | secondary | no | 2 | yes | yes | |
| 3 | 47 | blue-collar | married | unknown | no | 1506 | yes | no | |
| 4 | 33 | unknown | single | unknown | no | 1 | no | no | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 45206 | 51 | technician | married | tertiary | no | 825 | no | no | |
| 45207 | 71 | retired | divorced | primary | no | 1729 | no | no | |
| 45208 | 72 | retired | married | secondary | no | 5715 | no | no | |
| 45209 | 57 | blue-collar | married | secondary | no | 668 | no | no | |
| 45210 | 37 | entrepreneur | married | secondary | no | 2971 | no | no | |

| | contact | day | month | duration | campaign | pdays | previous | poutcome | y |
|-------|-----------|-----|-------|----------|----------|-------|----------|----------|-----|
| 0 | unknown | 5 | may | 261 | 1 | -1 | 0 | unknown | no |
| 1 | unknown | 5 | may | 151 | 1 | -1 | 0 | unknown | no |
| 2 | unknown | 5 | may | 76 | 1 | -1 | 0 | unknown | no |
| 3 | unknown | 5 | may | 92 | 1 | -1 | 0 | unknown | no |
| 4 | unknown | 5 | may | 198 | 1 | -1 | 0 | unknown | no |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 45206 | cellular | 17 | nov | 977 | 3 | -1 | 0 | unknown | yes |
| 45207 | cellular | 17 | nov | 456 | 2 | -1 | 0 | unknown | yes |
| 45208 | cellular | 17 | nov | 1127 | 5 | 184 | 3 | success | yes |
| 45209 | telephone | 17 | nov | 508 | 4 | -1 | 0 | unknown | no |
| 45210 | cellular | 17 | nov | 361 | 2 | 188 | 11 | other | no |

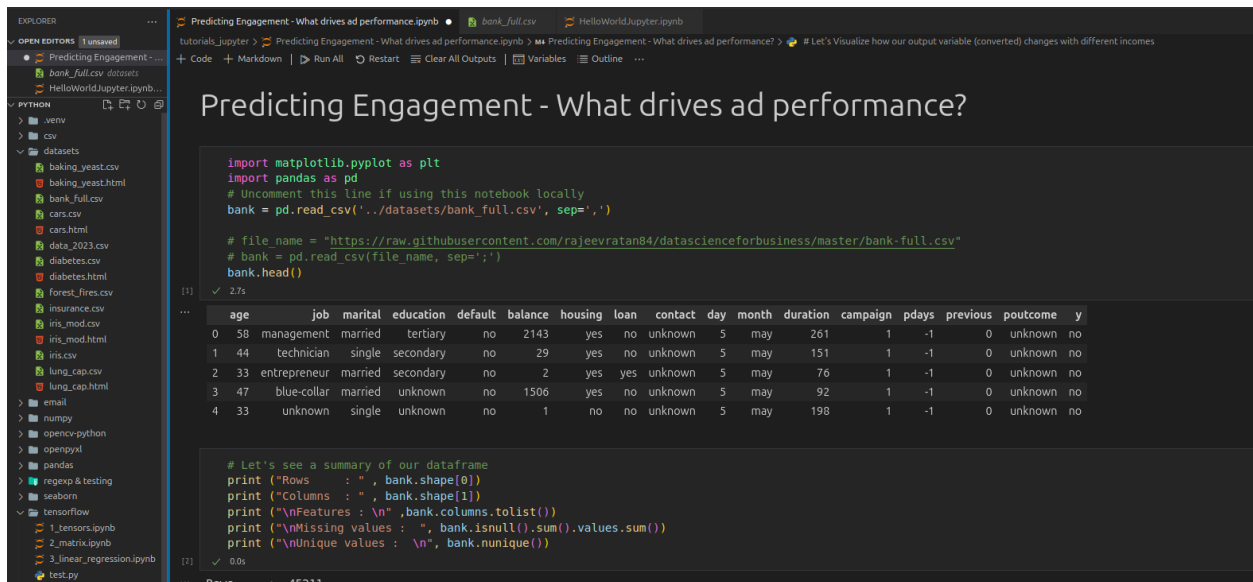
Run some Python code

Images above

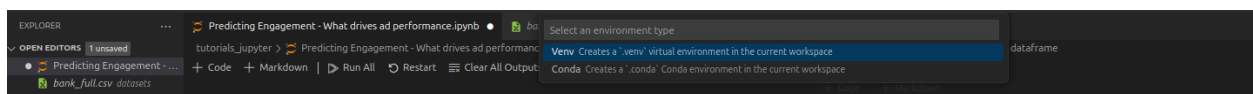
Self-Exploration

Another way of running Python – IDE

1. Is it possible to run a jupyter notebook that is integrated into Visual Studio Code via an extension Python language support.

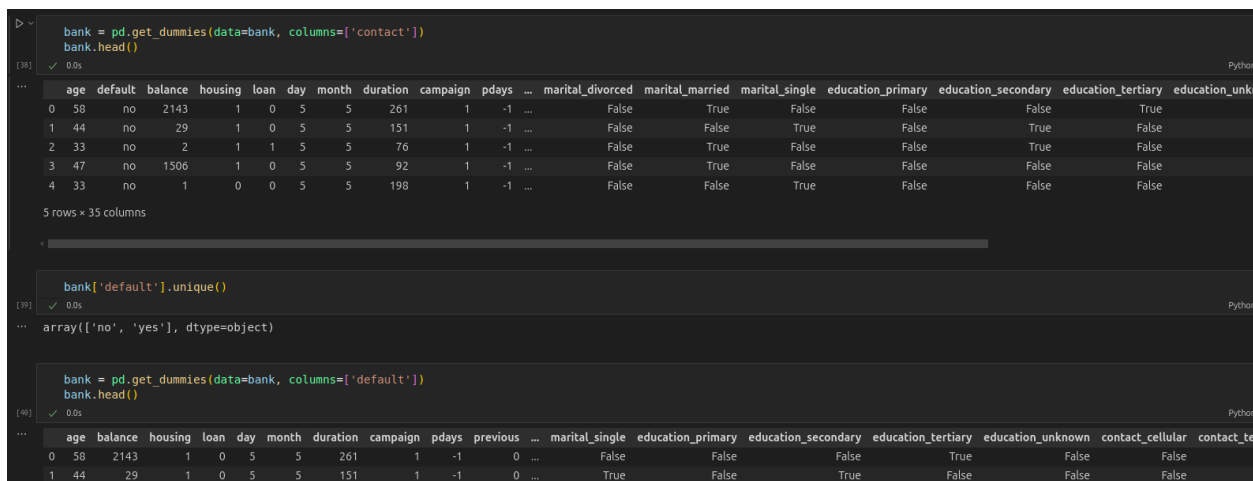


2. The extension allows creation of Python virtual environments with both conda and venv

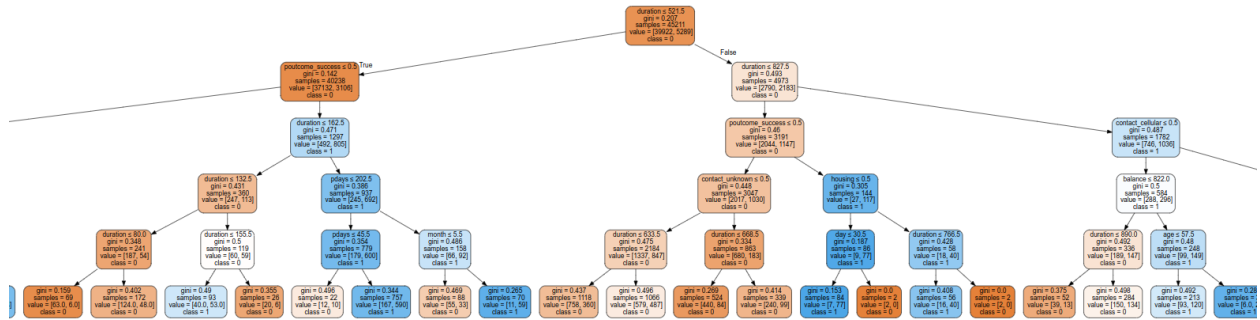


Predicting engagement

1. Run the script one by one to explore the data provided and transform wherever necessary.



2. Train a decision tree model with the refined dataset



Predicting Insurance Premiums

Predicting Insurance Premiums

- Our simple dataset contains a few attributes for each person such as
- Age, Sex, BMI, Children, Smoker, Region and their charges

Aim

- To use this info to predict charges for new customers

Code

Markdown

```

import pandas as pd

# Uncomment this line if using this notebook locally
insurance = pd.read_csv('../datasets/insurance.csv')

# file_name = "https://raw.githubusercontent.com/rajeevratan84/datascienceforbusiness/master/insurance.csv"
# insurance = pd.read_csv(file_name)

# Preview our data
insurance.head()

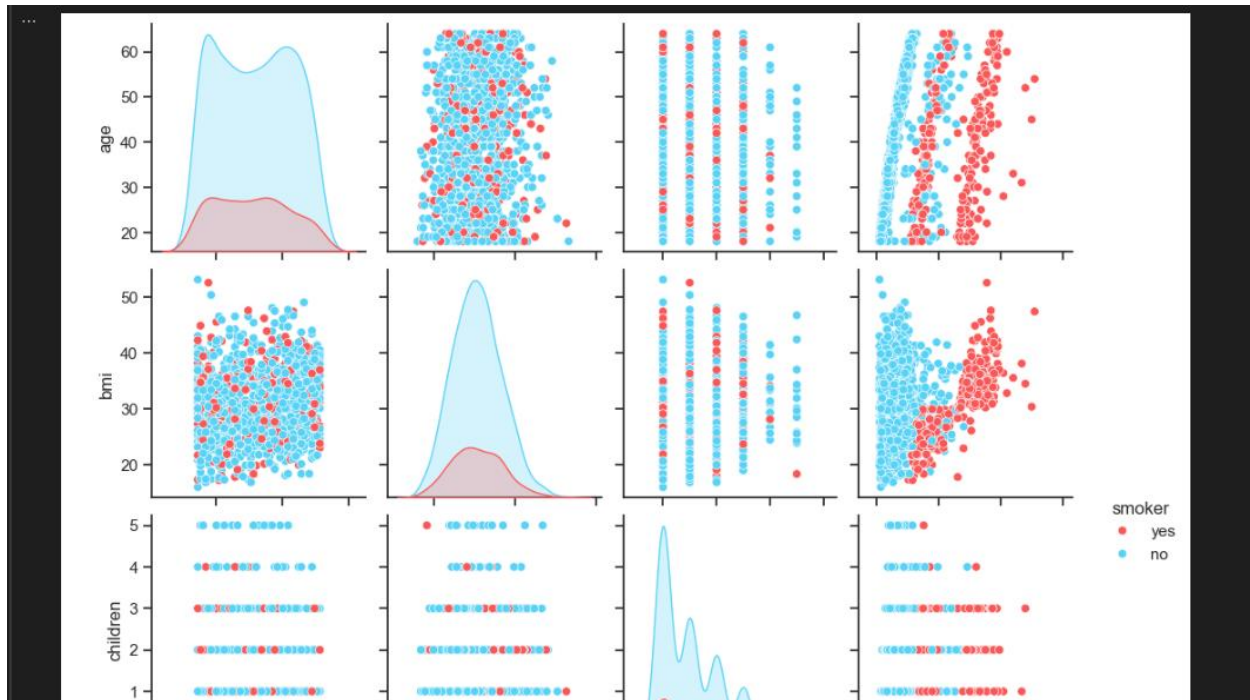
```

0.0s

Python

| | age | sex | bmi | children | smoker | region | charges |
|---|-----|--------|--------|----------|--------|-----------|-------------|
| 0 | 19 | female | 27.900 | 0 | yes | southwest | 16894.92400 |
| 1 | 18 | male | 33.770 | 1 | no | southeast | 1725.55230 |
| 2 | 28 | male | 33.000 | 3 | no | southeast | 4449.46200 |

Observe data features



Train models of different machine learning algorithms and observe effectiveness of each

```
# compare all results in one table
training_accuracies = [accuracy_MLR_train, accuracy_PR_train, accuracy_DTR_train, accuracy_RFR_train, accuracy_SVR_train]
testing_accuracies = [accuracy_MLR_test, accuracy_PR_test, accuracy_DTR_test, accuracy_RFR_test, accuracy_SVR_test]
training_RMSE = [RMSE_MLR_train, RMSE_PR_train, RMSE_DTR_train, RMSE_RFR_train, RMSE_SVR_train]
testing_RMSE = [RMSE_MLR_test, RMSE_PR_test, RMSE_DTR_test, RMSE_RFR_test, RMSE_SVR_test]
cv_accuracies = [accuracy_cv_MLR, accuracy_cv_PR, accuracy_cv_DTR, accuracy_cv_RFR, accuracy_cv_SVR]

parameters = ["fit_intercept=False", "fit_intercept=False", "max_depth=5", "n_estimators=400, max_depth=5", "kernel='li

table_data = {"Parameters": parameters, "Training Accuracy": training_accuracies, "Testing Accuracy": testing_accuracies,
              "Training RMSE": training_RMSE, "Testing RMSE": testing_RMSE, "10-Fold Score": cv_accuracies}
model_names = ["Multiple Linear Regression", "Polynomial Regression", "Decision Tree Regression", "Random Forest Regres

table_dataframe = pd.DataFrame(data=table_data, index=model_names)
table_dataframe
```

[46] ✓ 0.0s

| | Parameters | Training Accuracy | Testing Accuracy | Training RMSE | Testing RMSE | 10-Fold Score |
|----------------------------|-------------------------------|-------------------|------------------|---------------|--------------|---------------|
| Multiple Linear Regression | fit_intercept=False | -0.489561 | -0.324110 | 14589.307283 | 14438.166279 | 0.717113 |
| Polynomial Regression | fit_intercept=False | 0.835502 | 0.880879 | 4848.265611 | 4330.562228 | 0.839107 |
| Decision Tree Regression | max_depth=5 | 0.869426 | 0.871194 | 4319.509663 | 4503.167202 | 0.849424 |
| Random Forest Regression | n_estimators=400, max_depth=5 | 0.878632 | 0.896906 | 4164.457267 | 4028.712787 | 0.857379 |
| Support Vector Regression | kernel='linear', C=1000 | 0.652218 | 0.734317 | 7049.511742 | 6467.427432 | 0.705813 |